

# Phase Three Project Submission

Please fill out:

- Student name: Victor Keya
- Student pace: part time
- Scheduled project review date/time:
- Instructor name: Noah Kandie
- Blog post URL:

## 1. Introduction

### Business problem

SyriaTel company wants to reduce customer churn, which refers to customers switching to a different service provider. Customer churn can be costly for SyriaTel in terms of lost revenue and acquisition costs for new customers. By identifying customers who are likely to churn, SyriaTel can take proactive measures to retain them, such as offering special promotions or personalized services.

### Objectives

1. Implement a predictive churn modeling solution
2. Address customer retention.
3. Reduce churn rates.
4. Improve customer satisfaction and profitability.

## 2. Importing libraries and Loading Data

```
import pandas as pd
import csv
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import scipy.stats as stats
from pandas.plotting import parallel_coordinates
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import train_test_split, GridSearchCV
```

```
from sklearn.metrics import accuracy_score, confusion_matrix,
classification_report
%matplotlib inline
```

## 2.1 Load CSV File

```
df = pd.read_csv(r"C:\Users\Victor Keya\Documents\Flatiron\
Phase3_Project\bigml_59c28831336c6604c800002a.csv")
df
```

	state	account length	area code	phone number	international plan
0	KS	128	415	382-4657	no
1	OH	107	415	371-7191	no
2	NJ	137	415	358-1921	no
3	OH	84	408	375-9999	yes
4	OK	75	415	330-6626	yes
...	...	...	...	...	...
3328	AZ	192	415	414-4276	no
3329	WV	68	415	370-3271	no
3330	RI	28	510	328-8230	no
3331	CT	184	510	364-6381	yes
3332	TN	74	415	400-4344	no

	voice mail plan	number vmail messages	total day minutes
0	yes	25	265.1
1	yes	26	161.6
2	no	0	243.4
3	no	0	299.4
4	no	0	166.7
...	...	...	...
3328	yes	36	156.2
3329	no	0	231.1
3330	no	0	180.8
3331	no	0	213.8
3332	yes	25	234.4

	total day calls	total day charge	...	total eve calls
0	110	45.07	...	99

1	123	27.47	...	103
2	114	41.38	...	110
3	71	50.90	...	88
4	113	28.34	...	122
...	...	...	...	...
3328	77	26.55	...	126
3329	57	39.29	...	55
3330	109	30.74	...	58
3331	105	36.35	...	84
3332	113	39.85	...	82

	total eve charge	total night minutes	total night calls	\
0	16.78	244.7	91	
1	16.62	254.4	103	
2	10.30	162.6	104	
3	5.26	196.9	89	
4	12.61	186.9	121	
...	...	...	...	
3328	18.32	279.1	83	
3329	13.04	191.3	123	
3330	24.55	191.9	91	
3331	13.57	139.2	137	
3332	22.60	241.4	77	

	total night charge	total intl minutes	total intl calls	\
0	11.01	10.0	3	
1	11.45	13.7	3	
2	7.32	12.2	5	
3	8.86	6.6	7	
4	8.41	10.1	3	
...	...	...	...	
3328	12.56	9.9	6	
3329	8.61	9.6	4	
3330	8.64	14.1	6	
3331	6.26	5.0	10	
3332	10.86	13.7	4	

	total intl charge	customer service calls	churn
0	2.70	1	False
1	3.70	1	False
2	3.29	0	False
3	1.78	2	False
4	2.73	3	False
...	...	...	...
3328	2.67	2	False
3329	2.59	3	False
3330	3.81	2	False
3331	1.35	2	False
3332	3.70	0	False

[3333 rows x 21 columns]

## 2.2 Limitations

1. Phone Number Column: The 'phone number' column is likely irrelevant for predictive modeling and should be removed or transformed.
2. Geographic Information: The 'state' column may introduce regional biases and might need to be encoded properly or analyzed for its impact on churn.
3. Account Length: While this could be a significant feature, without context (e.g., is a longer account length generally better or worse?), it might be hard to interpret its impact directly.

## 3. Data Understanding

df.head()

	state	account length	area code	phone number	international	plan	\
0	KS	128	415	382-4657		no	
1	OH	107	415	371-7191		no	
2	NJ	137	415	358-1921		no	
3	OH	84	408	375-9999		yes	
4	OK	75	415	330-6626		yes	

	voice mail plan	number	vmail messages	total day minutes	total day	\
0	yes		25	265.1		
110						
1	yes		26	161.6		
123						
2	no		0	243.4		
114						
3	no		0	299.4		
71						
4	no		0	166.7		
113						

	total day charge	...	total eve calls	total eve charge	\
0	45.07	...	99	16.78	
1	27.47	...	103	16.62	
2	41.38	...	110	10.30	
3	50.90	...	88	5.26	
4	28.34	...	122	12.61	

	total night minutes	total night calls	total night charge	\
0	244.7	91	11.01	

1	254.4	103	11.45
2	162.6	104	7.32
3	196.9	89	8.86
4	186.9	121	8.41

	total intl minutes	total intl calls	total intl charge \
0	10.0	3	2.70
1	13.7	3	3.70
2	12.2	5	3.29
3	6.6	7	1.78
4	10.1	3	2.73

	customer service calls	churn
0	1	False
1	1	False
2	0	False
3	2	False
4	3	False

[5 rows x 21 columns]

df.tail()

	state	account length	area code	phone number	international plan
3328	AZ	192	415	414-4276	no
3329	WV	68	415	370-3271	no
3330	RI	28	510	328-8230	no
3331	CT	184	510	364-6381	yes
3332	TN	74	415	400-4344	no

	voice mail plan	number vmail messages	total day minutes \
3328	yes	36	156.2
3329	no	0	231.1
3330	no	0	180.8
3331	no	0	213.8
3332	yes	25	234.4

	total day calls	total day charge	...	total eve calls \
3328	77	26.55	...	126
3329	57	39.29	...	55
3330	109	30.74	...	58
3331	105	36.35	...	84
3332	113	39.85	...	82

	total eve charge	total night minutes	total night calls \
--	------------------	---------------------	---------------------

3328	18.32	279.1	83
3329	13.04	191.3	123
3330	24.55	191.9	91
3331	13.57	139.2	137
3332	22.60	241.4	77

	total night charge	total intl minutes	total intl calls \
3328	12.56	9.9	6
3329	8.61	9.6	4
3330	8.64	14.1	6
3331	6.26	5.0	10
3332	10.86	13.7	4

	total intl charge	customer service calls	churn
3328	2.67	2	False
3329	2.59	3	False
3330	3.81	2	False
3331	1.35	2	False
3332	3.70	0	False

[5 rows x 21 columns]

df.sample()

	state	account length	area code	phone number	international plan \
1570	NE	112	415	388-4282	no

	voice mail plan	number vmail messages	total day minutes \
1570	no	0	167.6

	total day calls	total day charge	...	total eve calls \
1570	100	28.49	...	90

	total eve charge	total night minutes	total night calls \
1570	13.13	281.4	107

	total night charge	total intl minutes	total intl calls \
1570	12.66	17.3	3

	total intl charge	customer service calls	churn
1570	4.67	2	False

[1 rows x 21 columns]

df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3333 entries, 0 to 3332
Data columns (total 21 columns):
```

#	Column	Non-Null Count	Dtype
0	state	3333 non-null	object
1	account length	3333 non-null	int64
2	area code	3333 non-null	int64
3	phone number	3333 non-null	object
4	international plan	3333 non-null	object
5	voice mail plan	3333 non-null	object
6	number vmail messages	3333 non-null	int64
7	total day minutes	3333 non-null	float64
8	total day calls	3333 non-null	int64
9	total day charge	3333 non-null	float64
10	total eve minutes	3333 non-null	float64
11	total eve calls	3333 non-null	int64
12	total eve charge	3333 non-null	float64
13	total night minutes	3333 non-null	float64
14	total night calls	3333 non-null	int64
15	total night charge	3333 non-null	float64
16	total intl minutes	3333 non-null	float64
17	total intl calls	3333 non-null	int64
18	total intl charge	3333 non-null	float64
19	customer service calls	3333 non-null	int64
20	churn	3333 non-null	bool

dtypes: bool(1), float64(8), int64(8), object(4)  
memory usage: 524.2+ KB

The dataset has 3333 rows and 21 columns.

No null values in the dataset.

Data types in the data set include floats, integers, objects and boolean.

```
df.describe()
```

	account length	area code	number vmail messages	total day minutes
count	3333.000000	3333.000000	3333.000000	3333.000000
mean	101.064806	437.182418	8.099010	179.775098
std	39.822106	42.371290	13.688365	54.467389
min	1.000000	408.000000	0.000000	0.000000
25%	74.000000	408.000000	0.000000	143.700000
50%	101.000000	415.000000	0.000000	179.400000
75%	127.000000	510.000000	20.000000	216.400000

max	243.000000	510.000000	51.000000
350.800000			
	total day calls	total day charge	total eve minutes
calls \			total eve
count	3333.000000	3333.000000	3333.000000
3333.000000			
mean	100.435644	30.562307	200.980348
100.114311			
std	20.069084	9.259435	50.713844
19.922625			
min	0.000000	0.000000	0.000000
0.000000			
25%	87.000000	24.430000	166.600000
87.000000			
50%	101.000000	30.500000	201.400000
100.000000			
75%	114.000000	36.790000	235.300000
114.000000			
max	165.000000	59.640000	363.700000
170.000000			
	total eve charge	total night minutes	total night calls \
count	3333.000000	3333.000000	3333.000000
mean	17.083540	200.872037	100.107711
std	4.310668	50.573847	19.568609
min	0.000000	23.200000	33.000000
25%	14.160000	167.000000	87.000000
50%	17.120000	201.200000	100.000000
75%	20.000000	235.300000	113.000000
max	30.910000	395.000000	175.000000
	total night charge	total intl minutes	total intl calls \
count	3333.000000	3333.000000	3333.000000
mean	9.039325	10.237294	4.479448
std	2.275873	2.791840	2.461214
min	1.040000	0.000000	0.000000
25%	7.520000	8.500000	3.000000
50%	9.050000	10.300000	4.000000
75%	10.590000	12.100000	6.000000
max	17.770000	20.000000	20.000000
	total intl charge	customer service calls	
count	3333.000000	3333.000000	
mean	2.764581	1.562856	
std	0.753773	1.315491	
min	0.000000	0.000000	
25%	2.300000	1.000000	
50%	2.780000	1.000000	



75%	3.270000	2.000000
max	5.400000	9.000000

This gives a summary of the distribution of the numeric data. From the count, we can see which columns have numeric data.

## 3.1 Data Cleaning

Lets make a copy of the dataset first

```
df1 = df.copy()
df1
```

	state	account length	area code	phone number	international plan
0	KS	128	415	382-4657	no
1	OH	107	415	371-7191	no
2	NJ	137	415	358-1921	no
3	OH	84	408	375-9999	yes
4	OK	75	415	330-6626	yes
...	...	...	...	...	...
3328	AZ	192	415	414-4276	no
3329	WV	68	415	370-3271	no
3330	RI	28	510	328-8230	no
3331	CT	184	510	364-6381	yes
3332	TN	74	415	400-4344	no

	voice mail plan	number vmail messages	total day minutes
0	yes	25	265.1
1	yes	26	161.6
2	no	0	243.4
3	no	0	299.4
4	no	0	166.7
...	...	...	...
3328	yes	36	156.2
3329	no	0	231.1
3330	no	0	180.8
3331	no	0	213.8
3332	yes	25	234.4

	total day calls	total day charge	...	total eve calls	\
0	110	45.07	...	99	
1	123	27.47	...	103	
2	114	41.38	...	110	
3	71	50.90	...	88	
4	113	28.34	...	122	
...	...	...	...	...	
3328	77	26.55	...	126	
3329	57	39.29	...	55	
3330	109	30.74	...	58	
3331	105	36.35	...	84	
3332	113	39.85	...	82	

	total eve charge	total night minutes	total night calls	\
0	16.78	244.7	91	
1	16.62	254.4	103	
2	10.30	162.6	104	
3	5.26	196.9	89	
4	12.61	186.9	121	
...	...	...	...	
3328	18.32	279.1	83	
3329	13.04	191.3	123	
3330	24.55	191.9	91	
3331	13.57	139.2	137	
3332	22.60	241.4	77	

	total night charge	total intl minutes	total intl calls	\
0	11.01	10.0	3	
1	11.45	13.7	3	
2	7.32	12.2	5	
3	8.86	6.6	7	
4	8.41	10.1	3	
...	...	...	...	
3328	12.56	9.9	6	
3329	8.61	9.6	4	
3330	8.64	14.1	6	
3331	6.26	5.0	10	
3332	10.86	13.7	4	

	total intl charge	customer service calls	churn
0	2.70	1	False
1	3.70	1	False
2	3.29	0	False
3	1.78	2	False
4	2.73	3	False
...	...	...	...
3328	2.67	2	False
3329	2.59	3	False
3330	3.81	2	False

```

3331          1.35          2  False
3332          3.70          0  False

[3333 rows x 21 columns]

```

Drop the Phone Number column

```
df1 = df1.drop('phone number', axis=1)
```

```
df1
```

```

      state  account length  area code international plan voice mail
plan \
0      KS          128      415          no
yes
1      OH          107      415          no
yes
2      NJ          137      415          no
no
3      OH           84      408          yes
no
4      OK           75      415          yes
no
...      ...      ...      ...      ...
..
3328    AZ          192      415          no
yes
3329    WV           68      415          no
no
3330    RI           28      510          no
no
3331    CT          184      510          yes
no
3332    TN           74      415          no
yes

      number vmail messages  total day minutes  total day calls \
0              25          265.1          110
1              26          161.6          123
2               0          243.4          114
3               0          299.4           71
4               0          166.7          113
...      ...      ...      ...
3328          36          156.2           77
3329           0          231.1           57
3330           0          180.8          109
3331           0          213.8          105
3332          25          234.4          113

      total day charge  total eve minutes  total eve calls  total eve

```

charge \			
0	45.07	197.4	99
16.78			
1	27.47	195.5	103
16.62			
2	41.38	121.2	110
10.30			
3	50.90	61.9	88
5.26			
4	28.34	148.3	122
12.61			
...	...	...	...
...			
3328	26.55	215.5	126
18.32			
3329	39.29	153.4	55
13.04			
3330	30.74	288.8	58
24.55			
3331	36.35	159.6	84
13.57			
3332	39.85	265.9	82
22.60			

	total night minutes	total night calls	total night charge \
0	244.7	91	11.01
1	254.4	103	11.45
2	162.6	104	7.32
3	196.9	89	8.86
4	186.9	121	8.41
...	...	...	...
3328	279.1	83	12.56
3329	191.3	123	8.61
3330	191.9	91	8.64
3331	139.2	137	6.26
3332	241.4	77	10.86

	total intl minutes	total intl calls	total intl charge \
0	10.0	3	2.70
1	13.7	3	3.70
2	12.2	5	3.29
3	6.6	7	1.78
4	10.1	3	2.73
...	...	...	...
3328	9.9	6	2.67
3329	9.6	4	2.59
3330	14.1	6	3.81
3331	5.0	10	1.35
3332	13.7	4	3.70

	customer service calls	churn
0	1	False
1	1	False
2	0	False
3	2	False
4	3	False
...	...	...
3328	2	False
3329	3	False
3330	2	False
3331	2	False
3332	0	False

[3333 rows x 20 columns]

Handle missing values

```
df1.isnull().sum()
```

state	0
account length	0
area code	0
international plan	0
voice mail plan	0
number vmail messages	0
total day minutes	0
total day calls	0
total day charge	0
total eve minutes	0
total eve calls	0
total eve charge	0
total night minutes	0
total night calls	0
total night charge	0
total intl minutes	0
total intl calls	0
total intl charge	0
customer service calls	0
churn	0

dtype: int64

Dataset has no missing values

Check for duplicates

```
df1.duplicated().sum()
```

0

No duplicates in this Dataset.

```
# Encode categorical variables
```

```
df1['international plan'] = df1['international plan'].map({'yes': 1, 'no': 0})
```

```
df1['voice mail plan'] = df1['voice mail plan'].map({'yes': 1, 'no': 0})
```

```
df1['churn'] = df1['churn'].astype(int)
```

```
df1
```

	state	account length	area code	international plan	voice mail plan \
--	-------	----------------	-----------	--------------------	-------------------

0	KS	128	415	0	
---	----	-----	-----	---	--

1					
---	--	--	--	--	--

1	OH	107	415	0	
---	----	-----	-----	---	--

1					
---	--	--	--	--	--

2	NJ	137	415	0	
---	----	-----	-----	---	--

0					
---	--	--	--	--	--

3	OH	84	408	1	
---	----	----	-----	---	--

0					
---	--	--	--	--	--

4	OK	75	415	1	
---	----	----	-----	---	--

0					
---	--	--	--	--	--

...	...	...	...	...	...
-----	-----	-----	-----	-----	-----

...					
-----	--	--	--	--	--

3328	AZ	192	415	0	
------	----	-----	-----	---	--

1					
---	--	--	--	--	--

3329	WV	68	415	0	
------	----	----	-----	---	--

0					
---	--	--	--	--	--

3330	RI	28	510	0	
------	----	----	-----	---	--

0					
---	--	--	--	--	--

3331	CT	184	510	1	
------	----	-----	-----	---	--

0					
---	--	--	--	--	--

3332	TN	74	415	0	
------	----	----	-----	---	--

1					
---	--	--	--	--	--

	number vmail messages	total day minutes	total day calls \
--	-----------------------	-------------------	-------------------

0	25	265.1	110
---	----	-------	-----

1	26	161.6	123
---	----	-------	-----

2	0	243.4	114
---	---	-------	-----

3	0	299.4	71
---	---	-------	----

4	0	166.7	113
---	---	-------	-----

...	...	...	...
-----	-----	-----	-----

3328	36	156.2	77
------	----	-------	----

3329	0	231.1	57
------	---	-------	----

3330	0	180.8	109
------	---	-------	-----

3331	0	213.8	105
------	---	-------	-----

3332	25	234.4	113
------	----	-------	-----

	total day charge	total eve minutes	total eve calls	total eve charge \
--	------------------	-------------------	-----------------	--------------------

0	45.07	197.4	99	
---	-------	-------	----	--

16.78			
1	27.47	195.5	103
16.62			
2	41.38	121.2	110
10.30			
3	50.90	61.9	88
5.26			
4	28.34	148.3	122
12.61			
...	...	...	...
...			
3328	26.55	215.5	126
18.32			
3329	39.29	153.4	55
13.04			
3330	30.74	288.8	58
24.55			
3331	36.35	159.6	84
13.57			
3332	39.85	265.9	82
22.60			
	total night minutes	total night calls	total night charge \
0	244.7	91	11.01
1	254.4	103	11.45
2	162.6	104	7.32
3	196.9	89	8.86
4	186.9	121	8.41
...	...	...	...
3328	279.1	83	12.56
3329	191.3	123	8.61
3330	191.9	91	8.64
3331	139.2	137	6.26
3332	241.4	77	10.86
	total intl minutes	total intl calls	total intl charge \
0	10.0	3	2.70
1	13.7	3	3.70
2	12.2	5	3.29
3	6.6	7	1.78
4	10.1	3	2.73
...	...	...	...
3328	9.9	6	2.67
3329	9.6	4	2.59
3330	14.1	6	3.81
3331	5.0	10	1.35
3332	13.7	4	3.70
	customer service calls	churn	
0	1	0	

1	1	0
2	0	0
3	2	0
4	3	0
...	...	...
3328	2	0
3329	3	0
3330	2	0
3331	2	0
3332	0	0

[3333 rows x 20 columns]

*# Calculate the churn rate*

```
churn_rate = df1['Churn'].mean()
```

*# Display the churn rate*

```
print(f'Overall Churn Rate: {churn_rate:.2%}')
```

Overall Churn Rate: 14.49%

*# Check data types*

```
print("\nData Types:")
```

```
print(df1.dtypes)
```

Data Types:

state	object
account length	int64
area code	int64
international plan	int64
voice mail plan	int64
number vmail messages	int64
total day minutes	float64
total day calls	int64
total day charge	float64
total eve minutes	float64
total eve calls	int64
total eve charge	float64
total night minutes	float64
total night calls	int64
total night charge	float64
total intl minutes	float64
total intl calls	int64
total intl charge	float64
customer service calls	int64
churn	int32
dtype:	object

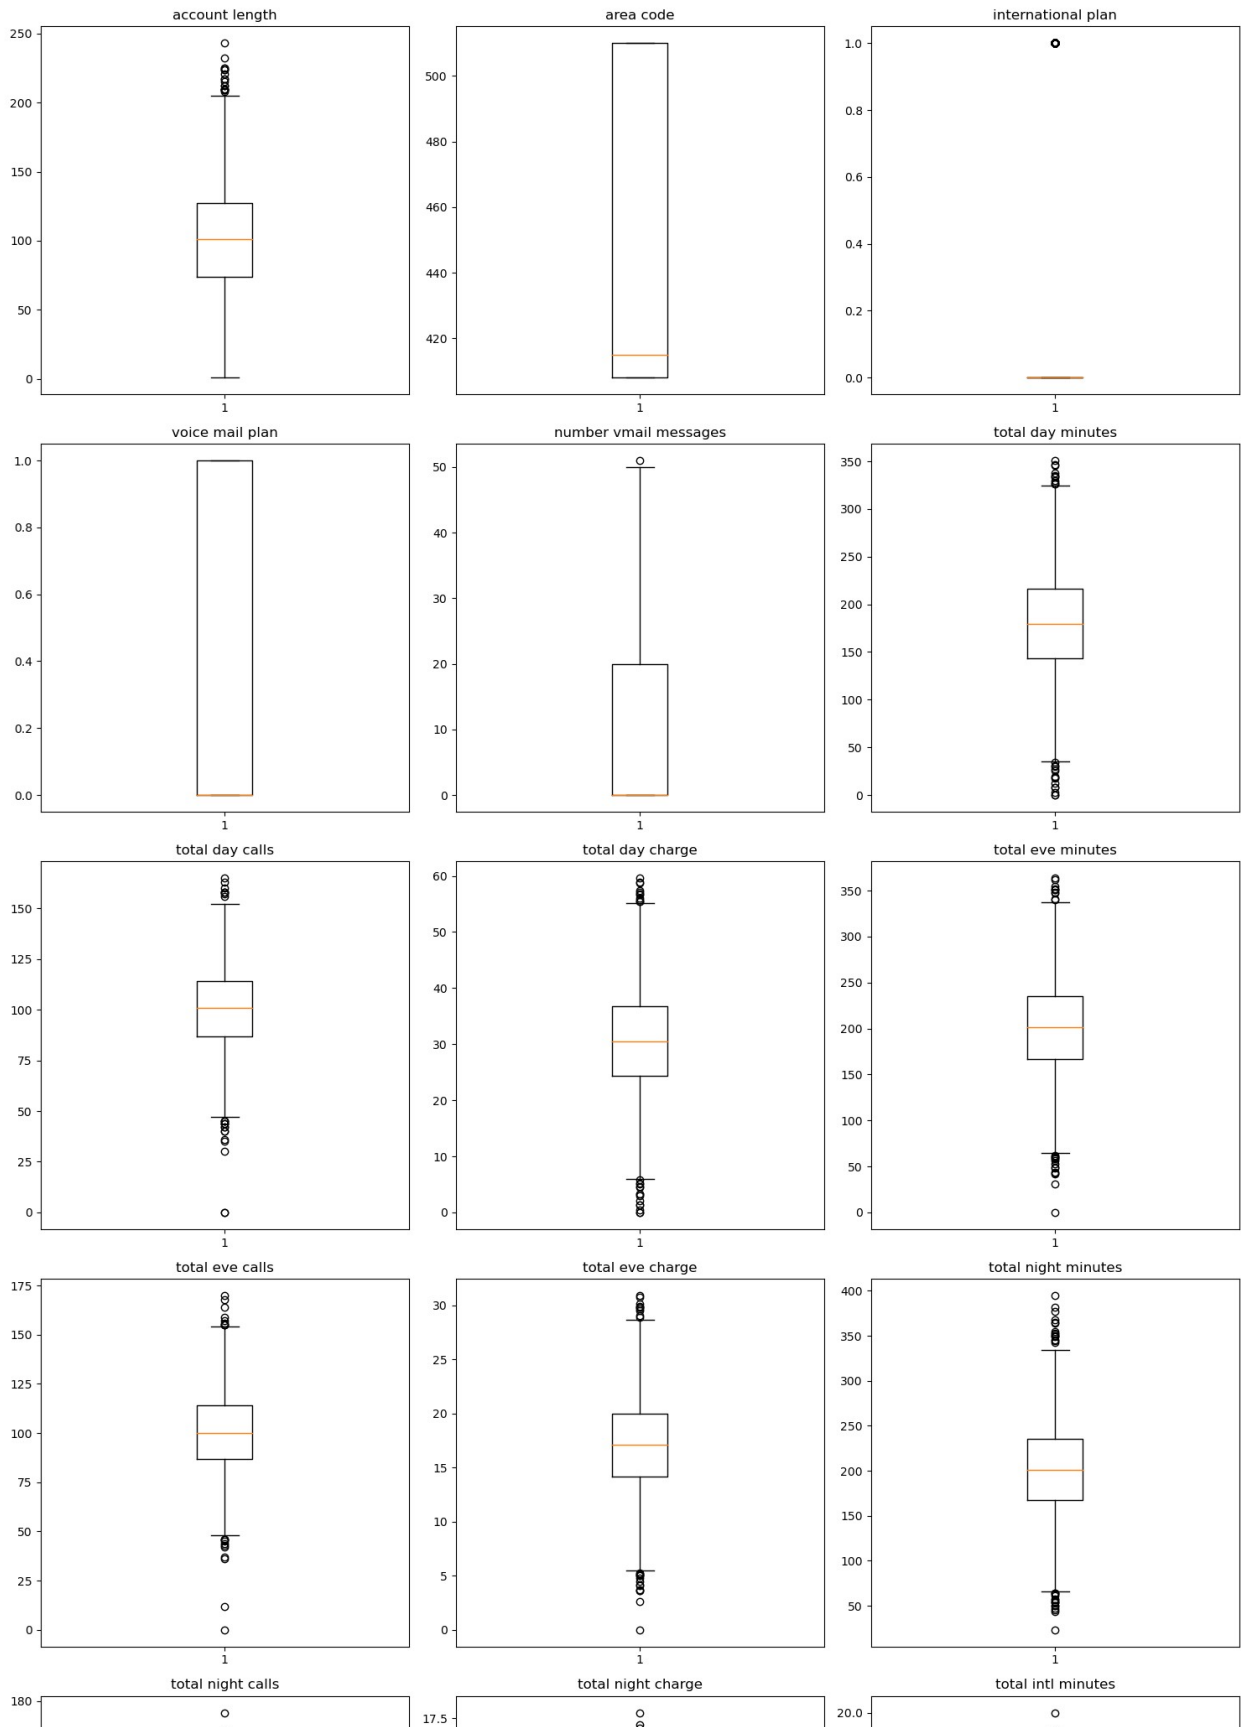


## Outliers

```
numeric_columns = df1.select_dtypes(include=['float64', 'int64',
'int32'])

# Plot box plots for each numeric column
num_cols = len(numeric_columns.columns)
cols_per_row = 3
num_rows = (num_cols - 1) // cols_per_row + 1

plt.figure(figsize=(15, 5 * num_rows))
for i, col in enumerate(numeric_columns.columns):
    plt.subplot(num_rows, cols_per_row, i+1)
    plt.boxplot(numeric_columns[col])
    plt.title(col)
plt.tight_layout()
plt.show()
```



## Skewness and Kurtosis

```
# Drop non-numeric columns to focus on numerical data
numeric_df = df1.drop(['state'], axis=1)

# Calculate skewness
skewness = numeric_df.skew()
print("Skewness:\n", skewness)

# Calculate kurtosis
kurtosis = numeric_df.kurt()
print("Kurtosis:\n", kurtosis)
```

Skewness:

account length	0.096606
area code	1.126823
international plan	2.726332
voice mail plan	0.999140
number vmail messages	1.264824
total day minutes	-0.029077
total day calls	-0.111787
total day charge	-0.029083
total eve minutes	-0.023877
total eve calls	-0.055563
total eve charge	-0.023858
total night minutes	0.008921
total night calls	0.032500
total night charge	0.008886
total intl minutes	-0.245136
total intl calls	1.321478
total intl charge	-0.245287
customer service calls	1.091359
churn	2.018356

dtype: float64

Kurtosis:

account length	-0.107836
area code	-0.705632
international plan	5.436146
voice mail plan	-1.002322
number vmail messages	-0.051129
total day minutes	-0.019940
total day calls	0.243182
total day charge	-0.019812
total eve minutes	0.025630
total eve calls	0.206156
total eve charge	0.025487
total night minutes	0.085816
total night calls	-0.072020
total night charge	0.085663
total intl minutes	0.609185
total intl calls	3.083589

```
total intl charge      0.609610
customer service calls 1.730914
churn                  2.075006
dtype: float64
```

## Skewness Interpretation

Highly Positively Skewed: international plan (2.73), churn (2.02) indicate that most customers do not have an international plan and do not churn, but a small number do.

Moderately Positively Skewed: area code (1.13), voice mail plan (1.00), number vmail messages (1.26), total intl calls (1.32), customer service calls (1.09).

Moderately Negatively Skewed: total intl minutes (-0.25), total intl charge (-0.25).

## Kurtosis Interpretation

Highly Leptokurtic: international plan (5.44) indicating a distribution with heavy tails, this suggests many values are far from the mean.

Moderately Leptokurtic: total intl calls (3.08), customer service calls (1.73), churn (2.08).

Near-Normal Kurtosis: Many features are close to 0, suggests that their distributions are not heavy-tailed.

## Business Insights:

The skewness in customer service calls and churn indicates most customers do not frequently contact customer service or churn, highlighting the need to focus on the minority who do.

High kurtosis in international plan and churn suggests significant differences between customers who churn and those who do not, indicating targeted strategies could be effective.

## Feature Engineering

```
# Feature engineering
df1['total_calls'] = df1['total day calls'] + df1['total eve calls'] +
df1['total night calls']
df1['total_minutes'] = df1['total day minutes'] + df1['total eve
minutes'] + df1['total night minutes']
df1['total_charge'] = df1['total day charge'] + df1['total eve
charge'] + df1['total night charge'] + df1['total intl charge']

# Interaction Features
df1['day_calls_per_minute'] = df1['total day calls'] / df1['total day
minutes']
df1['eve_calls_per_minute'] = df1['total eve calls'] / df1['total eve
minutes']
df1['night_calls_per_minute'] = df1['total night calls'] / df1['total
night minutes']
```

```
# Drop non-numeric and irrelevant columns
```

```
df1.drop(['state', 'area code'], axis=1, inplace=True)
```

```
df1
```

	account length	international plan	voice mail plan	\
0	128	0	1	
1	107	0	1	
2	137	0	0	
3	84	1	0	
4	75	1	0	
...	...	...	...	
3328	192	0	1	
3329	68	0	0	
3330	28	0	0	
3331	184	1	0	
3332	74	0	1	

	number vmail messages	total day minutes	total day calls	\
0	25	265.1	110	
1	26	161.6	123	
2	0	243.4	114	
3	0	299.4	71	
4	0	166.7	113	
...	...	...	...	
3328	36	156.2	77	
3329	0	231.1	57	
3330	0	180.8	109	
3331	0	213.8	105	
3332	25	234.4	113	

	total day charge	total eve minutes	total eve calls	total eve charge \
0	45.07	197.4	99	
16.78				
1	27.47	195.5	103	
16.62				
2	41.38	121.2	110	
10.30				
3	50.90	61.9	88	
5.26				
4	28.34	148.3	122	
12.61				
...	...	...	...	
...				
3328	26.55	215.5	126	
18.32				
3329	39.29	153.4	55	
13.04				
3330	30.74	288.8	58	

24.55				
3331	36.35	159.6	84	
13.57				
3332	39.85	265.9	82	
22.60				
	...	total intl calls	total intl charge	customer service calls
churn \				
0	...	3	2.70	1
0				
1	...	3	3.70	1
0				
2	...	5	3.29	0
0				
3	...	7	1.78	2
0				
4	...	3	2.73	3
0				
...	...	...	...	...
...				
3328	...	6	2.67	2
0				
3329	...	4	2.59	3
0				
3330	...	6	3.81	2
0				
3331	...	10	1.35	2
0				
3332	...	4	3.70	0
0				
	total_calls	total_minutes	total_charge	
day_calls_per_minute \				
0	300	707.2	75.56	0.414938
1	329	611.5	59.24	0.761139
2	328	527.2	62.29	0.468365
3	248	558.2	66.80	0.237141
4	356	501.9	52.09	0.677864
...	...	...	...	...
3328	286	650.8	60.10	0.492958
3329	235	575.8	63.53	0.246646
3330	258	661.5	67.74	0.602876

3331	326	512.6	57.53	0.491113
3332	272	741.7	77.01	0.482082

	eve_calls_per_minute	night_calls_per_minute
0	0.501520	0.371884
1	0.526854	0.404874
2	0.907591	0.639606
3	1.421648	0.452006
4	0.822657	0.647405
...	...	...
3328	0.584687	0.297384
3329	0.358540	0.642969
3330	0.200831	0.474205
3331	0.526316	0.984195
3332	0.308387	0.318973

[3333 rows x 24 columns]

df1.shape

(3333, 24)

*#Lets capitalise the column titles.*

df1.columns = [col.capitalize() for col in df1.columns]

df1.columns

Index(['Account length', 'International plan', 'Voice mail plan',  
'Number vmil messages', 'Total day minutes', 'Total day  
calls',  
'Total day charge', 'Total eve minutes', 'Total eve calls',  
'Total eve charge', 'Total night minutes', 'Total night calls',  
'Total night charge', 'Total intl minutes', 'Total intl calls',  
'Total intl charge', 'Customer service calls', 'Churn',  
'Total\_calls',  
'Total\_minutes', 'Total\_charge', 'Day\_calls\_per\_minute',  
'Eve\_calls\_per\_minute', 'Night\_calls\_per\_minute'],  
dtype='object')

df1.head()

	Account length	International plan	Voice mail plan	Number vmil messages \
0	128	0	1	
25				
1	107	0	1	
26				
2	137	0	0	
0				

3	84	1	0	
0				
4	75	1	0	
0				
Total day minutes    Total day calls    Total day charge    Total eve				
minutes \				
0	265.1	110	45.07	
197.4				
1	161.6	123	27.47	
195.5				
2	243.4	114	41.38	
121.2				
3	299.4	71	50.90	
61.9				
4	166.7	113	28.34	
148.3				
Total eve calls    Total eve charge    ...    Total intl calls \				
0	99	16.78	3	
1	103	16.62	3	
2	110	10.30	5	
3	88	5.26	7	
4	122	12.61	3	
Total intl charge    Customer service calls    Churn    Total_calls \				
0	2.70	1	0	300
1	3.70	1	0	329
2	3.29	0	0	328
3	1.78	2	0	248
4	2.73	3	0	356
Total_minutes    Total_charge    Day_calls_per_minute				
Eve_calls_per_minute \				
0	707.2	75.56	0.414938	
0.501520				
1	611.5	59.24	0.761139	
0.526854				
2	527.2	62.29	0.468365	
0.907591				
3	558.2	66.80	0.237141	
1.421648				
4	501.9	52.09	0.677864	
0.822657				
Night_calls_per_minute				
0	0.371884			
1	0.404874			
2	0.639606			
3	0.452006			



```
4                0.647405
```

```
[5 rows x 24 columns]
```

## 4. Exploratory Data Analysis

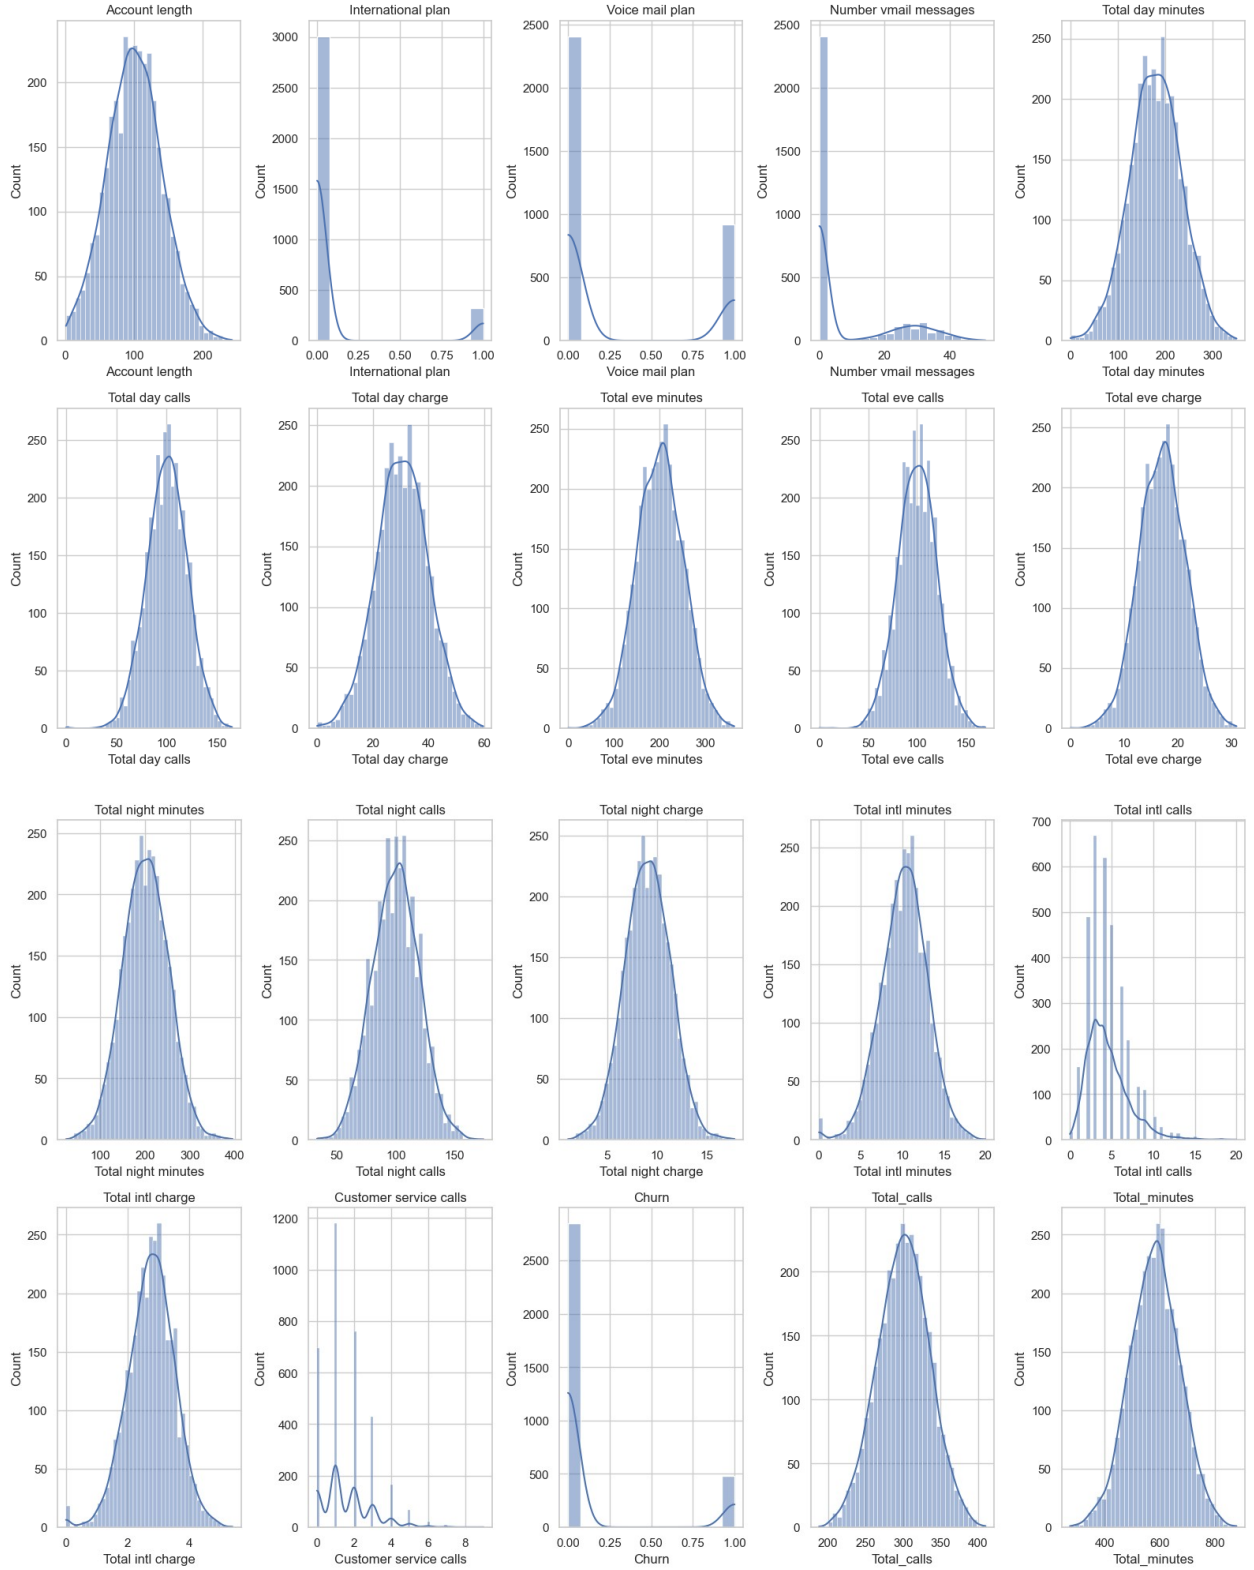
### 4.1. Univariate Analysis


```
np.random.seed(0)

# 2. Distribution Plots using Seaborn
sns.set(style="whitegrid")

# Plot histograms and KDE for numerical features
numerical_features = dfl.select_dtypes(include=[np.number]).columns

# Split the features into chunks of 10 to avoid the subplot limit
error
chunk_size = 10
for i in range(0, len(numerical_features), chunk_size):
    plt.figure(figsize=(16, 10))
    for j, feature in enumerate(numerical_features[i:i + chunk_size]):
        plt.subplot(2, 5, j + 1)
        sns.histplot(dfl[feature], kde=True)
        plt.title(feature)
    plt.tight_layout()
    plt.show()
```





## Histogram Analysis

### Account length:

This graph shows a nearly normal distribution centered around 100 days, indicating that most customers have had their accounts for about 100 days, with fewer customers having very short or very long account durations.

### International plan:

This histogram is highly skewed to the left, showing that the vast majority of customers (near 0 on the x-axis) do not have an international plan. A small number of customers (near 1 on the x-axis) have opted for the international plan.

### Voice mail plan:

Similar to the international plan, this distribution is also heavily skewed to the left. Most customers do not have a voice mail plan (0 on the x-axis), while a minority have the plan (1 on the x-axis).

### Number vmail messages:

This distribution is heavily right-skewed, indicating that most customers have not used the voicemail service (0 messages), and only a few customers have left a significant number of voicemail messages.


### Total day minutes:

This graph is approximately normally distributed, centered around 200 minutes. It shows that most customers use around 200 minutes of daytime calling, with fewer customers having significantly lower or higher usage.

## 4.2 Bivariate Analysis

```
# Correlation Heatmap
plt.figure(figsize=(16, 10))
```

```
correlation_matrix = df1.corr()
sns.heatmap(correlation_matrix, annot=True, fmt=".2f",
cmap='coolwarm', linewidths=0.5)
plt.title('Correlation Heatmap')
plt.show()
```



Correlation coefficients range from -1 to 1, where:

1 indicates a perfect positive correlation, -1 indicates a perfect negative correlation, 0 indicates no correlation.

High Positive Correlations:

Total day minutes, Total day charge: These have a correlation coefficient of 1.00, indicating a perfect positive correlation. This makes sense as the charge is directly proportional to the usage. Total eve minutes, Total eve charge: Similar to the day minutes and charge, this pair also shows a perfect positive correlation. Total night minutes, Total night charge: This pair follows the same pattern with a perfect positive correlation. Total intl minutes, Total intl charge: This pair also has a perfect positive correlation.

Negative Correlations:

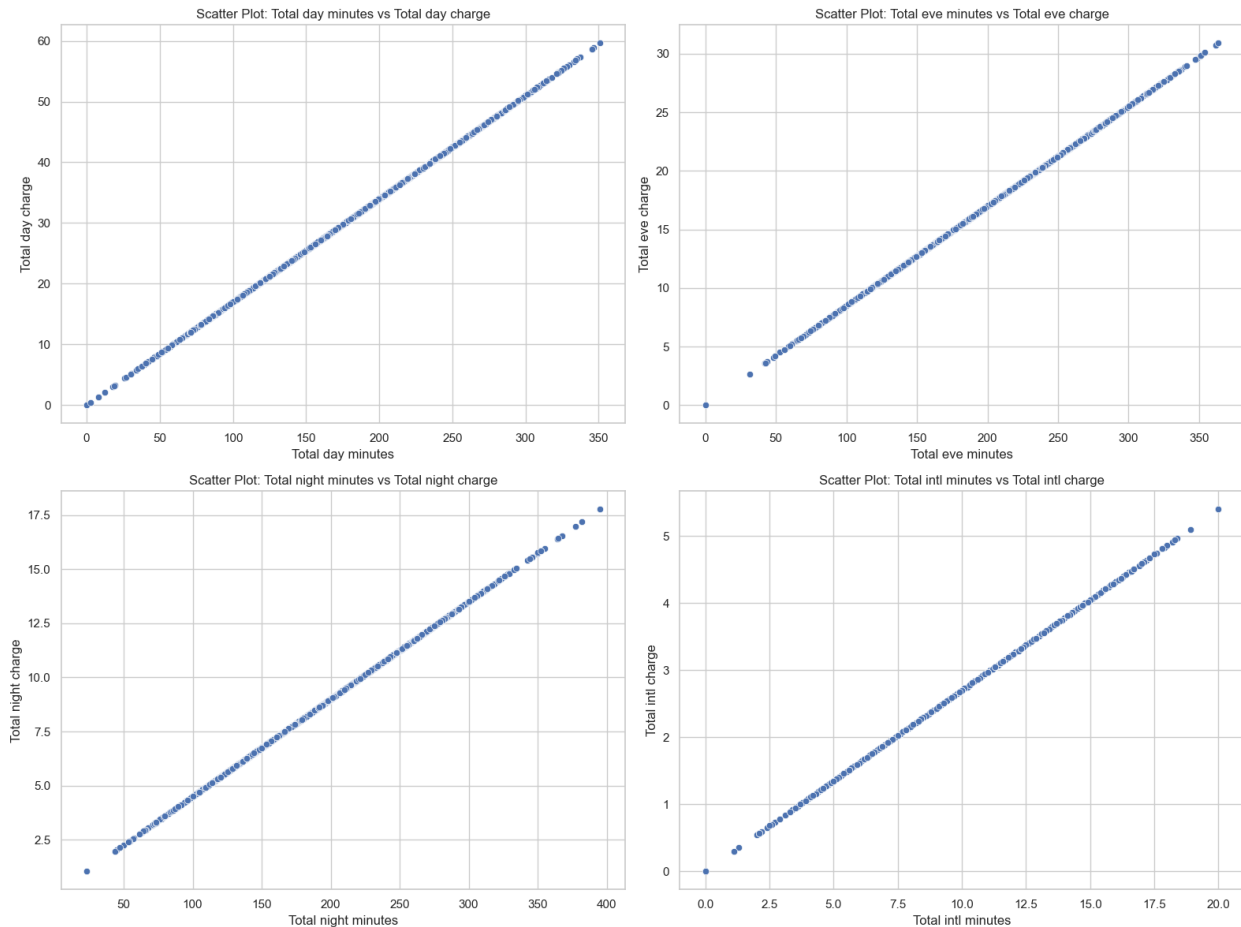
Day\_calls\_per\_minute and Total day minutes: There is a moderate negative correlation (-0.37), suggesting that as the number of total day minutes increases, the number of calls per minute decreases. Eve\_calls\_per\_minute and Total eve minutes: There is a moderate negative correlation (-0.49). Night\_calls\_per\_minute and Total night minutes: There is a moderate negative correlation (-0.74), suggesting that as the total night minutes increase, the number of night calls per minute decreases significantly.

Low or No Correlations:

Account length and most other variables: The correlations here are very low, indicating that the length of time a customer has had an account does not significantly correlate with their usage patterns or likelihood to churn. Customer service calls and most other variables: These generally show low correlations, except for churn, where there is a moderate positive correlation (0.21), indicating that customers who call customer service more frequently are somewhat more likely to churn.

```
# Scatter plots for selected pairs of variables
pairs = [
    ('Total day minutes', 'Total day charge'),
    ('Total eve minutes', 'Total eve charge'),
    ('Total night minutes', 'Total night charge'),
    ('Total intl minutes', 'Total intl charge')
]

plt.figure(figsize=(16, 12))
for i, (var1, var2) in enumerate(pairs):
    plt.subplot(2, 2, i + 1)
    sns.scatterplot(data=df1, x=var1, y=var2)
    plt.title(f'Scatter Plot: {var1} vs {var2}')
plt.tight_layout()
plt.show()
```



These scatter plots display the relationship between the number of minutes used and the corresponding charges across different time periods (day, evening, night, and international).

#### 1. Total Day Minutes vs. Total Day Charge:

The scatter plot shows a perfect linear relationship, indicating that as the total day minutes increase, the total day charge increases proportionally. This strong positive linear correlation confirms that the billing is directly proportional to the usage of day minutes, which is expected as charges are calculated based on usage.

#### 1. Total Eve Minutes vs. Total Eve Charge:

Similar to the day minutes and charges, this scatter plot also shows a perfect linear relationship. The total evening minutes are directly proportional to the total evening charge, indicating a consistent billing structure for evening calls.

#### 1. Total Night Minutes vs. Total Night Charge:

This plot also exhibits a perfect linear relationship, demonstrating that as the total night minutes increase, the total night charge increases in a directly proportional manner. This consistency across different times of the day (day, evening, night) suggests a straightforward billing system where charges are directly based on the minutes used.

#### 1. Total Intl Minutes vs. Total Intl Charge:

The scatter plot of total international minutes vs. total international charge also shows a perfect linear relationship. This indicates that international charges are directly proportional to the international minutes used, consistent with the patterns observed in domestic call charges.

All scatter plots exhibit a perfect positive linear relationship between minutes used and corresponding charges, reflecting a proportional billing system. This suggests that the telecom company's charging structure is consistent and linear, meaning customers are charged a fixed rate per minute for their calls, regardless of whether the calls are made during the day, evening, night, or internationally.

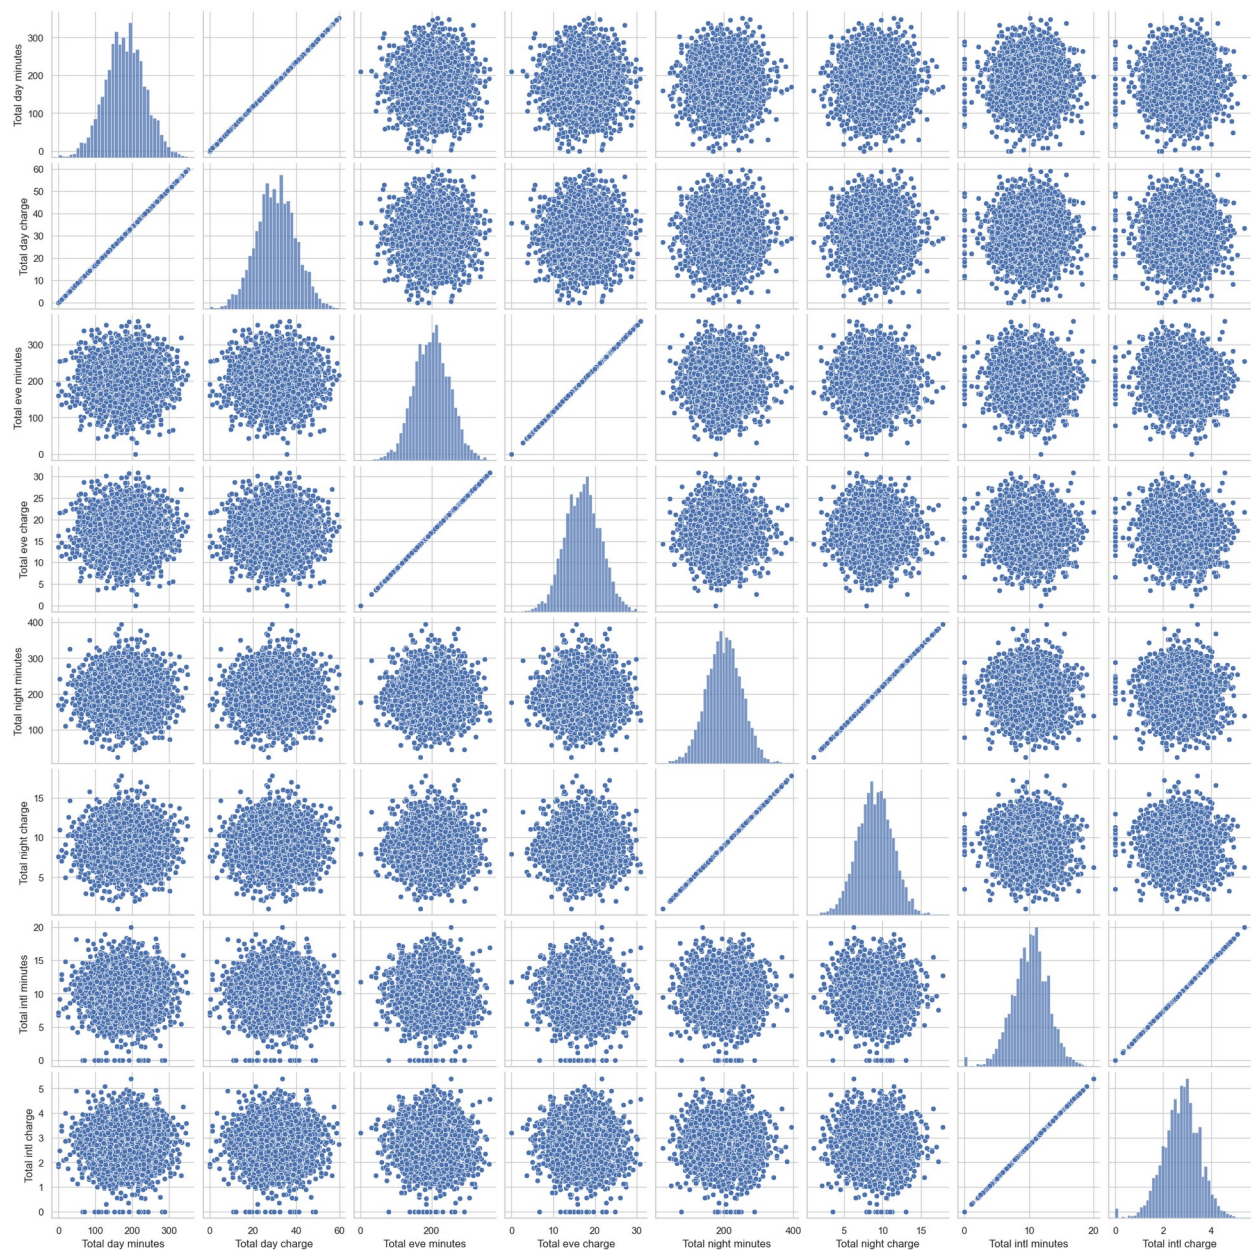
```
# Pair plots for a subset of the features
subset_features = [
    'Total day minutes', 'Total day charge',
    'Total eve minutes', 'Total eve charge',
    'Total night minutes', 'Total night charge',
    'Total intl minutes', 'Total intl charge'
]

sns.pairplot(df1[subset_features])
plt.suptitle('Pair Plot of Selected Features', y=1.02)
plt.show()

C:\Users\Victor Keya\Documents\Anaconda\Lib\site-packages\seaborn\
axisgrid.py:118: UserWarning: The figure layout has changed to tight
self._figure.tight_layout(*args, **kwargs)
```



Pair Plot of Selected Features



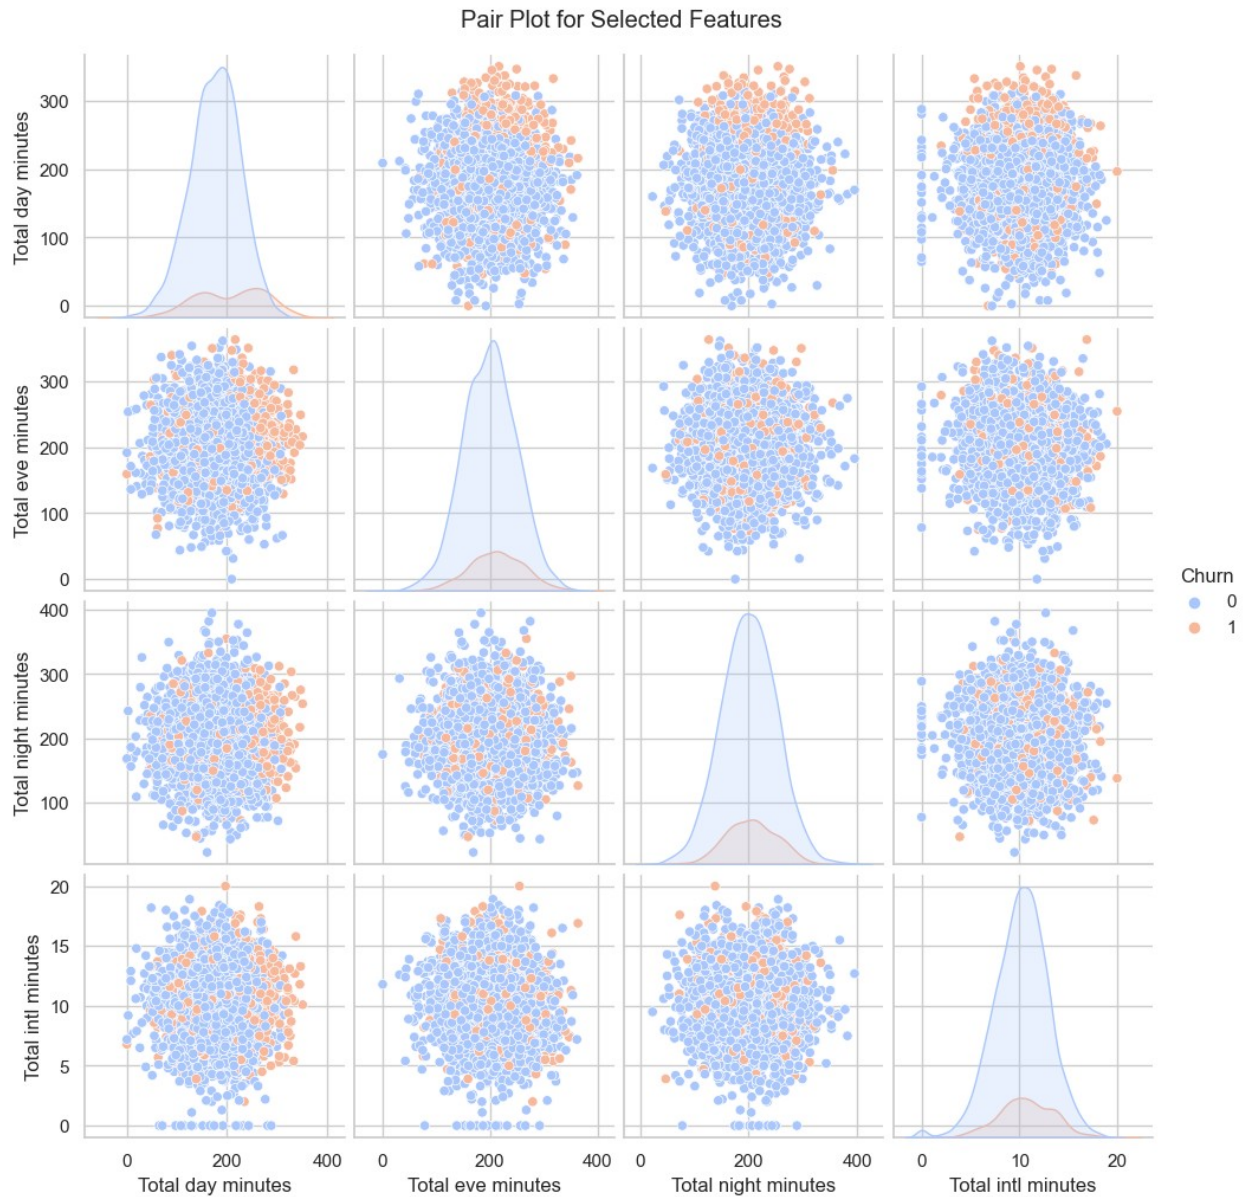
## 4.3 Multivariate Analysis

```
# Pair plot for a subset of features
plt.figure(figsize=(10, 10))
subset_features = ['Total day minutes', 'Total eve minutes', 'Total
night minutes', 'Total intl minutes', 'Churn']
sns.pairplot(df1[subset_features], hue='Churn', palette='coolwarm')
plt.suptitle('Pair Plot for Selected Features', y=1.02)
plt.show()
```



```
C:\Users\Victor Keya\Documents\Anaconda\Lib\site-packages\seaborn\axisgrid.py:118: UserWarning: The figure layout has changed to tight
self._figure.tight_layout(*args, **kwargs)
```

<Figure size 1000x1000 with 0 Axes>



This pair plot visualizes the relationships between selected features in the SyriaTel dataset, specifically focusing on various types of minutes used by customers and their churn status. Each dot represents a customer, with blue indicating non-churned customers (Churn = 0) and orange indicating churned customers (Churn = 1).

Key Observations:

Diagonal Plots:

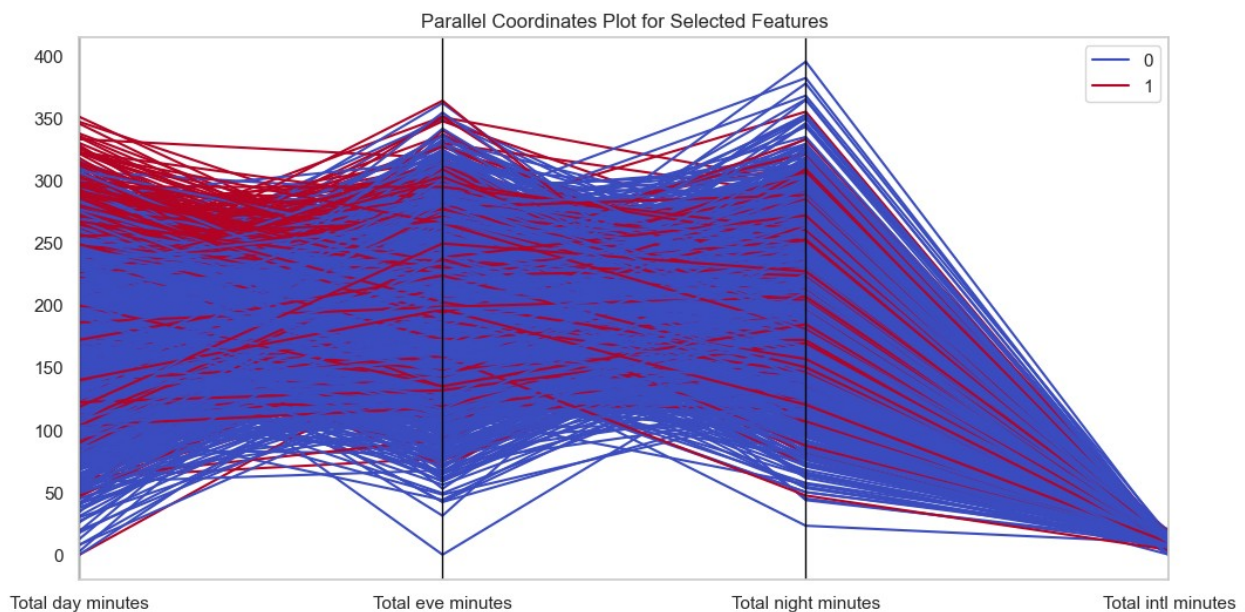
These plots show the distribution of each feature. The kernel density estimation (KDE) curves show the distribution for churned (orange) and non-churned (blue) customers. Total Day Minutes, Total Eve Minutes, Total Night Minutes, Total Intl Minutes: The distributions for churned and non-churned customers are similar, with a slight difference in the intensity of orange and blue, indicating that churned and non-churned customers have somewhat similar usage patterns across these time periods. Off-Diagonal Plots:

These plots show the scatter relationships between pairs of features. The scatter plots do not show strong linear relationships between the features, suggesting that individual minute usage features might not strongly predict each other.

Churn Patterns: Churn vs. Non-Churn: There is a mix of blue and orange dots in most scatter plots, indicating that churned customers have usage patterns that are broadly similar to those of non-churned customers. Slight differences in distribution might be more apparent with advanced statistical methods but are not immediately obvious in this pair plot visualization.

Insights: Overall Similarity: The usage patterns of churned and non-churned customers are largely similar across the different types of call minutes. Feature Independence: The lack of strong visible correlations between different types of call minutes suggests that these features might act independently of each other in influencing churn.

```
# Parallel coordinates plot for a subset of features
plt.figure(figsize=(12, 6))
parallel_features = ['Total day minutes', 'Total eve minutes', 'Total night minutes', 'Total intl minutes', 'Churn']
parallel_coordinates(df1[parallel_features], class_column='Churn', colormap='coolwarm')
plt.title('Parallel Coordinates Plot for Selected Features')
plt.show()
```



This parallel coordinates plot visualizes the relationships between selected features in the SyriaTel dataset and the churn status of customers. Each line represents a customer, with blue lines indicating non-churned customers (Churn = 0) and red lines indicating churned customers (Churn = 1)

## 5. Data Preprocessing

### 5.1 Train-Test Split

```
#Handle Missing Values
df1.fillna(df1.mean(), inplace=True)

# Define features and target
X = df1.drop(columns=['Churn'])
y = df1['Churn']

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=1)
```

### 5.2 Feature Scaling

```
# Standardize the features
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
X_test_scaled

array([[ -0.81666076,  -0.32942967,  -0.61058295, ...,  -0.3729483 ,
        -0.12880568,   0.54110596],
       [  0.02926645,  -0.32942967,  -0.61058295, ...,   0.52194962,
        -1.01010079,   1.26341822],
       [ -0.71713991,  -0.32942967,  -0.61058295, ...,  -0.74120513,
         2.74310286,  -0.44815204],
       ...,
       [  0.52687068,  -0.32942967,  -0.61058295, ...,  -0.35641865,
        -0.04224218,  -0.10615051],
       [  0.75079259,  -0.32942967,  -0.61058295, ...,  -0.13244234,
        -0.16340447,   0.35462517],
       [  0.12878729,  -0.32942967,  -0.61058295, ...,  -0.22918336,
        -0.83882258,   1.74494651]])
```

### 5.3 Model Training

#### 5.3.1 Logistic Regression Model

```
# Train logistic regression model
model = LogisticRegression(random_state=1)
model.fit(X_train_scaled, y_train)
```

```

LogisticRegression(random_state=1)

# Make predictions on the test set
y_pred = model.predict(X_test_scaled)

# Calculate accuracy
accuracy = accuracy_score(y_test, y_pred)
print(f'Accuracy: {accuracy}')

# Display the confusion matrix
conf_matrix = confusion_matrix(y_test, y_pred)
print('Confusion Matrix:')
print(conf_matrix)

# Display the classification report
class_report = classification_report(y_test, y_pred)
print('Classification Report:')
print(class_report)

# Display model coefficients
coefficients = pd.DataFrame(model.coef_[0], X.columns,
                             columns=['Coefficient'])
coefficients['Odds Ratio'] = np.exp(coefficients['Coefficient'])
print(coefficients)

```

Accuracy: 0.8455772113943029

Confusion Matrix:

```
[[548  21]
 [ 82  16]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.87	0.96	0.91	569
1	0.43	0.16	0.24	98
accuracy			0.85	667
macro avg	0.65	0.56	0.58	667
weighted avg	0.81	0.85	0.81	667

	Coefficient	Odds Ratio
Account length	0.029762	1.030210
International plan	0.621476	1.861674
Voice mail plan	-0.767956	0.463960
Number vmail messages	0.362316	1.436653
Total day minutes	0.248454	1.282041
Total day calls	-0.041988	0.958881
Total day charge	0.248824	1.282516
Total eve minutes	0.089268	1.093374
Total eve calls	0.008920	1.008960
Total eve charge	0.090330	1.094535

Total night minutes	-0.071785	0.930731
Total night calls	0.130843	1.139789
Total night charge	-0.073361	0.929265
Total intl minutes	0.115986	1.122980
Total intl calls	-0.294597	0.744832
Total intl charge	0.115763	1.122729
Customer service calls	0.711644	2.037338
Total_calls	0.055147	1.056696
Total_minutes	0.160331	1.173899
Total_charge	0.248694	1.282349
Day_calls_per_minute	0.242937	1.274988
Eve_calls_per_minute	0.040695	1.041534
Night_calls_per_minute	-0.337875	0.713284

## Logistic Regression Model Justification

Logistic Regression is simple and interpretable.

It shows clear coefficients and odds ratios, making it easy to understand the impact of each feature on churn.

However, it struggles with class imbalance, as indicated by low precision and recall for the minority class (churn).

## Logistic Regression Results

### Business Insights:

1. **Understanding Customer Churn High Precision for Non-Churners:** The model accurately predicts customers who are not likely to churn (precision of 0.87 and recall of 0.96). This means the company can trust the model when it says a customer is not likely to churn. **Low Precision and Recall for Churners:** The model struggles to accurately predict customers who are likely to churn (precision of 0.43 and recall of 0.16). This indicates that the model is not very reliable in identifying customers who will churn, potentially missing many at-risk customers.
2. **Key Drivers of Churn Positive Association with Churn: International Plan:** Customers with an international plan are more likely to churn. The company might need to investigate if the international plan's value proposition is clear and competitive. **Total Day Minutes/Charge:** Higher usage during the day is associated with churn. This might indicate dissatisfaction with peak-time service quality or pricing. **Customer Service Calls:** Customers who contact customer service frequently are more likely to churn. This is a critical insight suggesting that customer service interactions need improvement.
3. **Negative Association with Churn: Voice Mail Plan:** Customers with a voice mail plan are less likely to churn. This might indicate satisfaction with additional services. **Total Night Minutes/Charge:** Higher usage during night times is associated with lower churn. This might suggest that services during off-peak times are satisfactory.

## Recommendations for Business Improvement

1. **Enhancing Customer Service Training and Quality:** Invest in customer service training and quality improvements to reduce the number of issues leading to customer service calls. **Proactive Support:** Implement proactive customer support to resolve issues before they escalate, based on frequent service calls.
2. **Product and Service Review International Plan Evaluation:** Review and possibly revamp the international plan. Conduct customer surveys to understand pain points and improve the offering. **Daytime Service Quality:** Investigate the quality of service during peak times. Improve network infrastructure if needed to enhance the customer experience during high usage periods.
3. **Customer Retention Strategies Targeted Promotions:** Use the model's predictions to target likely churners with personalized promotions and incentives to stay. **Loyalty Programs:** Develop loyalty programs for customers with high usage, especially those using the service during nights and having voice mail plans.
4. **Feature Improvements Voice Mail Plan:** Promote voice mail plans as they are associated with lower churn. Consider bundling voice mail plans with other services to increase overall customer satisfaction.
5. **Data-Driven Decision Making Continuous Monitoring:** Regularly update the model with new data to keep track of changing customer behaviors and improve prediction accuracy over time. **Expand Features:** Collect more data on customer behavior and feedback to enhance the model's predictive power and provide deeper insights.

### 5.3.2 Decision Tree Model

```
# Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.3, random_state=1)

# Initialize and train the Decision Tree model
model = DecisionTreeClassifier(random_state=1)
model.fit(X_train, y_train)

# Predict on the test set
y_pred = model.predict(X_test)

# Evaluate the model
accuracy = accuracy_score(y_test, y_pred)
report = classification_report(y_test, y_pred)

print(f'Accuracy: {accuracy}')
print(f'Classification Report:\n{report}')
```

Accuracy: 0.943

Classification Report:

	precision	recall	f1-score	support
0	0.97	0.96	0.97	855
1	0.79	0.83	0.81	145

accuracy			0.94	1000
macro avg	0.88	0.90	0.89	1000
weighted avg	0.94	0.94	0.94	1000

## Decision Tree Model Justification

Decision Tree provides a high accuracy and a good performance for both classes.

It captures complex patterns in the data but can overfit easily.

Less interpretable compared to Logistic Regression but provides feature importance.

## Business Insights

### 1. Key Drivers of Churn:

Check if features such as "Customer service calls" and "Total day minutes" are significant predictors:

**Customer Service Calls:** High number of customer service calls might indicate unresolved issues. The company can improve customer service quality by training agents better, reducing wait times, and resolving issues more effectively.

**Total Day Minutes:** High usage might be correlated with dissatisfaction if it leads to high charges. Consider introducing more competitive pricing plans or offering discounts to high-usage customers.

### 1. Targeted Retention Campaigns: The company can segment customers based on their likelihood of churning.

**High-Risk Customers:** Implement targeted retention strategies such as personalized offers, discounts, loyalty programs, and proactive customer service outreach to high-risk customers.

**Medium-Risk Customers:** Monitor these customers and offer periodic incentives to ensure they remain satisfied.

### 1. Resource Allocation: Efficiently allocate resources to areas that will have the most impact on reducing churn.

**Customer Support:** Allocate more resources to improve customer service for customers identified as high-risk. **Marketing Efforts:** Direct marketing efforts towards high-risk customers with personalized engagement strategies.

### 1. Pricing Strategies: If the model indicates that high charges are a significant predictor of churn, the company can reconsider its pricing strategy.

**Flexible Pricing Plans:** Introduce more flexible or usage-based pricing plans to cater to different customer segments.



## Implementing Insights

### Feature Importance Analysis:

```
#Identify and list the top features influencing churn.

feature_importances = model.feature_importances_
features = X.columns
feature_importance_df = pd.DataFrame({'Feature': features,
'Importance': feature_importances})
feature_importance_df =
feature_importance_df.sort_values(by='Importance', ascending=False)
print(feature_importance_df)
```

	Feature	Importance
19	Total_charge	0.424844
16	Customer service calls	0.129130
3	Number vmail messages	0.113011
15	Total intl charge	0.078476
1	International plan	0.071625
14	Total intl calls	0.069650
18	Total_minutes	0.025016
5	Total day calls	0.019376
8	Total eve calls	0.013838
0	Account length	0.010808
11	Total night calls	0.008297
10	Total night minutes	0.007544
12	Total night charge	0.005305
2	Voice mail plan	0.004325
21	Eve_calls_per_minute	0.003594
6	Total day charge	0.003460
20	Day_calls_per_minute	0.003432
9	Total eve charge	0.002956
22	Night_calls_per_minute	0.002595
4	Total day minutes	0.002307
13	Total intl minutes	0.000411
7	Total eve minutes	0.000000
17	Total_calls	0.000000

### Prediction and Customer Segmentation:

```
#Predict churn probabilities
churn_probabilities = model.predict_proba(X_test_scaled)[: , 1] #
Assuming class 1 is churn

C:\Users\Victor Keya\Documents\Anaconda\Lib\site-packages\sklearn\
base.py:464: UserWarning: X does not have valid feature names, but
DecisionTreeClassifier was fitted with feature names
  warnings.warn(
```



### 5.3.3 Random Forest Classifier

```
# Split the data
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

# Scaling
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Model Selection and Hyperparameter Tuning
#param_grid = {
    #'n_estimators': [50, 100, 200],
    #'max_depth': [None, 10, 20],
    #'min_samples_split': [2, 5, 10],
    #'min_samples_leaf': [1, 2, 4]
#}
#rf = RandomForestClassifier(random_state=42)
#grid_search = GridSearchCV(estimator=rf, param_grid=param_grid, cv=5,
scoring='accuracy')
#grid_search.fit(X_train, y_train)

#best_params = grid_search.best_params_
#best_model = grid_search.best_estimator_

#Train the Random Forest model
model = RandomForestClassifier(random_state=1)
model.fit(X_train_scaled, y_train)

# Predict on the test set
y_pred = model.predict(X_test_scaled)

# Calculate accuracy
accuracy = accuracy_score(y_test, y_pred)
print(f'Accuracy: {accuracy}')
```

```
# Print classification report
report = classification_report(y_test, y_pred)
print(report)

# Print confusion matrix
conf_matrix = confusion_matrix(y_test, y_pred)
print(conf_matrix)

# Feature importances
feature_importances = model.feature_importances_
features = X.columns
feature_importance_df = pd.DataFrame({'Feature': features,
'Importance': feature_importances})
```

```
feature_importance_df =
feature_importance_df.sort_values(by='Importance', ascending=False)
print(feature_importance_df)
```

Accuracy: 0.9805097451274363

	precision	recall	f1-score	support
0	0.98	1.00	0.99	566
1	1.00	0.87	0.93	101
accuracy			0.98	667
macro avg	0.99	0.94	0.96	667
weighted avg	0.98	0.98	0.98	667

```
[[566  0]
 [ 13 88]]
```

	Feature	Importance
19	Total_charge	0.177486
16	Customer service calls	0.124370
1	International plan	0.069990
18	Total_minutes	0.065448
4	Total day minutes	0.065063
6	Total day charge	0.065059
2	Voice mail plan	0.035000
3	Number vmail messages	0.034710
14	Total intl calls	0.034375
13	Total intl minutes	0.033676
20	Day_calls_per_minute	0.031905
15	Total intl charge	0.031098
7	Total eve minutes	0.027055
9	Total eve charge	0.026072
5	Total day calls	0.022722
21	Eve_calls_per_minute	0.022664
10	Total night minutes	0.020773
22	Night_calls_per_minute	0.019771
11	Total night calls	0.019446
0	Account length	0.018767
8	Total eve calls	0.018517
12	Total night charge	0.018179
17	Total_calls	0.017855

## Random Forest Classifier Model Justification

Random Forest provides the highest accuracy and balanced performance across both classes.

It is less prone to overfitting compared to a single Decision Tree.

Provides a robust understanding of feature importance.

## Business Insights and Actions

1. Address High Charges:

Analyze the relationship between high total charges and churn. Consider offering discounts or loyalty programs to high-usage customers to enhance their satisfaction and reduce churn risk.

#### 1. Improve Customer Service:

Since frequent customer service calls are a strong churn predictor, focus on improving the quality of customer service. Implement training programs for service representatives and ensure quick resolution of customer issues.

#### 1. Monitor High Usage:

Keep an eye on customers with high total day minutes and total minutes. Provide them with personalized offers or services to increase their satisfaction.

#### 1. International Plan Customers:

Pay special attention to customers with international plans. Ensure their specific needs are met, possibly by providing better international call rates or packages.

### Which Model Is best for My DataSet

Random Forest emerges as the best model based on:

Accuracy: Highest among the three models (0.97).

Balanced Performance: High precision and recall for both churn and non-churn classes, indicating robustness in handling class imbalance.

Feature Importance: Provides clear insights into the most influential factors affecting churn, which can guide strategic business decisions.

Generalization: Less likely to overfit compared to Decision Trees, making it more reliable for unseen data.

df1				
	Account length	International plan	Voice mail plan	\
0	128	0	1	
1	107	0	1	
2	137	0	0	
3	84	1	0	
4	75	1	0	
...	...	...	...	
3328	192	0	1	
3329	68	0	0	
3330	28	0	0	
3331	184	1	0	
3332	74	0	1	
	Number vmail messages	Total day minutes	Total day calls	\
0	25	265.1	110	
1	26	161.6	123	
2	0	243.4	114	

3		0	299.4	71
4		0	166.7	113
...	...	...	...	...
3328		36	156.2	77
3329		0	231.1	57
3330		0	180.8	109
3331		0	213.8	105
3332		25	234.4	113
Total day charge    Total eve minutes    Total eve calls    Total eve				
charge \				
0	45.07	197.4	99	
16.78				
1	27.47	195.5	103	
16.62				
2	41.38	121.2	110	
10.30				
3	50.90	61.9	88	
5.26				
4	28.34	148.3	122	
12.61				
...	...	...	...	
...				
3328	26.55	215.5	126	
18.32				
3329	39.29	153.4	55	
13.04				
3330	30.74	288.8	58	
24.55				
3331	36.35	159.6	84	
13.57				
3332	39.85	265.9	82	
22.60				
...    Total intl calls    Total intl charge    Customer service calls				
Churn \				
0	...	3	2.70	1
0				
1	...	3	3.70	1
0				
2	...	5	3.29	0
0				
3	...	7	1.78	2
0				
4	...	3	2.73	3
0				
...	...	...	...	...
...				
3328	...	6	2.67	2

```

0
3329 ... 4 2.59 3
0
3330 ... 6 3.81 2
0
3331 ... 10 1.35 2
0
3332 ... 4 3.70 0
0

```

```

      Total_calls Total_minutes Total_charge
Day_calls_per_minute \
0      300      707.2      75.56      0.414938
1      329      611.5      59.24      0.761139
2      328      527.2      62.29      0.468365
3      248      558.2      66.80      0.237141
4      356      501.9      52.09      0.677864
...      ...      ...      ...      ...
3328      286      650.8      60.10      0.492958
3329      235      575.8      63.53      0.246646
3330      258      661.5      67.74      0.602876
3331      326      512.6      57.53      0.491113
3332      272      741.7      77.01      0.482082

```

```

      Eve_calls_per_minute Night_calls_per_minute
0      0.501520      0.371884
1      0.526854      0.404874
2      0.907591      0.639606
3      1.421648      0.452006
4      0.822657      0.647405
...      ...      ...
3328      0.584687      0.297384
3329      0.358540      0.642969
3330      0.200831      0.474205
3331      0.526316      0.984195
3332      0.308387      0.318973

```

```
[3333 rows x 24 columns]
```